


```
In [1]: import tensorflow as tf
import matplotlib.pyplot as plt
from tensorflow import keras
import numpy as np
```

```
In [2]: (x_train, y_train), (x_test, y_test) = keras.datasets.fashion_mnist.load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz>
29515/29515 [=====] - 0s 0us/step
Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz>
26421880/26421880 [=====] - 0s 0us/step
Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz>
5148/5148 [=====] - 0s 0us/step
Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz>
4422102/4422102 [=====] - 0s 0us/step


```
In [4]: plt.imshow(x_train[1])
```

```
Out[4]: <matplotlib.image.AxesImage at 0x7f134effcf40>
```



```
In [5]: plt.imshow(x_train[0])
```

```
Out[5]: <matplotlib.image.AxesImage at 0x7f13492692e0>
```



```
In [6]: x_train = x_train.astype('float32') / 255.0  
x_test = x_test.astype('float32') / 255.0  
x_train = x_train.reshape(-1, 28, 28, 1)  
x_test = x_test.reshape(-1, 28, 28, 1)
```

```
In [8]: x_train.shape
```

```
Out[8]: (60000, 28, 28, 1)
```

```
In [9]: x_test.shape
```

```
Out[9]: (10000, 28, 28, 1)
```

```
In [10]: y_train.shape
```

```
Out[10]: (60000,)
```

```
In [11]: y_test.shape
```

```
Out[11]: (10000,)
```

```
In [12]: model = keras.Sequential([  
    keras.layers.Conv2D(32, (3,3), activation='relu', input_shape=(28,28,1)),  
    keras.layers.MaxPooling2D((2,2)),  
    keras.layers.Dropout(0.25),  
    keras.layers.Conv2D(64, (3,3), activation='relu'),  
    keras.layers.MaxPooling2D((2,2)),  
    keras.layers.Dropout(0.25),  
    keras.layers.Conv2D(128, (3,3), activation='relu'),  
    keras.layers.Flatten(),  
    keras.layers.Dense(128, activation='relu'),  
])
```

```
keras.layers.Dropout(0.25),
keras.layers.Dense(10, activation='softmax')
])
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
dropout (Dropout)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
dropout_1 (Dropout)	(None, 5, 5, 64)	0
conv2d_2 (Conv2D)	(None, 3, 3, 128)	73856
flatten (Flatten)	(None, 1152)	0
dense (Dense)	(None, 128)	147584
dropout_2 (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 10)	1290
=====		
Total params: 241,546		
Trainable params: 241,546		
Non-trainable params: 0		

```
In [13]: model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accu
history = model.fit(x_train, y_train, epochs=10, validation_data=(x_test, y_test))
```

```

Epoch 1/10
1875/1875 [=====] - 79s 41ms/step - loss: 0.5641 - accuracy:
0.7937 - val_loss: 0.3837 - val_accuracy: 0.8573
Epoch 2/10
1875/1875 [=====] - 76s 41ms/step - loss: 0.3717 - accuracy:
0.8630 - val_loss: 0.3210 - val_accuracy: 0.8839
Epoch 3/10
1875/1875 [=====] - 78s 42ms/step - loss: 0.3274 - accuracy:
0.8800 - val_loss: 0.2994 - val_accuracy: 0.8894
Epoch 4/10
1875/1875 [=====] - 75s 40ms/step - loss: 0.3020 - accuracy:
0.8880 - val_loss: 0.3065 - val_accuracy: 0.8839
Epoch 5/10
1875/1875 [=====] - 78s 42ms/step - loss: 0.2834 - accuracy:
0.8956 - val_loss: 0.2740 - val_accuracy: 0.8997
Epoch 6/10
1875/1875 [=====] - 76s 41ms/step - loss: 0.2694 - accuracy:
0.8994 - val_loss: 0.2656 - val_accuracy: 0.8984
Epoch 7/10
1875/1875 [=====] - 77s 41ms/step - loss: 0.2597 - accuracy:
0.9026 - val_loss: 0.2845 - val_accuracy: 0.9002
Epoch 8/10
1875/1875 [=====] - 77s 41ms/step - loss: 0.2511 - accuracy:
0.9066 - val_loss: 0.2661 - val_accuracy: 0.9016
Epoch 9/10
1875/1875 [=====] - 77s 41ms/step - loss: 0.2393 - accuracy:
0.9097 - val_loss: 0.2755 - val_accuracy: 0.9022
Epoch 10/10
1875/1875 [=====] - 76s 40ms/step - loss: 0.2352 - accuracy:
0.9108 - val_loss: 0.2523 - val_accuracy: 0.9066

```

```
In [15]: test_loss, test_acc = model.evaluate(x_test, y_test)
```

```

313/313 [=====] - 4s 13ms/step - loss: 0.2523 - accuracy: 0.
9066

```

```
In [16]: print('Test accuracy:', test_acc)
```

```
Test accuracy: 0.9065999984741211
```