

```
!pip install git+https://github.com/andreinechaev/nvcc4jupyter.git
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Collecting git+https://github.com/andreinechaev/nvcc4jupyter.git
  Cloning https://github.com/andreinechaev/nvcc4jupyter.git to /tmp/pip-req-build-52_yj781
  Running command git clone --filter=blob:none --quiet https://github.com/andreinechaev/nvcc4jupyter.git /tmp/pip-req-build-52_yj781
  Resolved https://github.com/andreinechaev/nvcc4jupyter.git to commit aac710a35f52bb78ab34d2e52517237941399eff
  Preparing metadata (setup.py) ... done
Building wheels for collected packages: NVCCPlugin
  Building wheel for NVCCPlugin (setup.py) ... done
  Created wheel for NVCCPlugin: filename=NVCCPlugin-0.0.2-py3-none-any.whl size=4305 sha256=cb139076a076c203ff73e44e4a1e24b60ca92eb75514
  Stored in directory: /tmp/pip-ephem-wheel-cache-o4q78u9x/wheels/db/c1/1f/a2bb07bbb4a1ce3c43921252aeafaa6205f08637e292496f04
Successfully built NVCCPlugin
Installing collected packages: NVCCPlugin
Successfully installed NVCCPlugin-0.0.2
```

```
%load_ext nvcc_plugin
```

```
created output directory at /content/src
Out bin /content/result.out
```

```
%%cu
```

```
/*
 * file name: matrix.cu
 *
 * matrix.cu contains the code that realize some common used matrix operations in CUDA
 *
 * this is a toy program for learning CUDA, some functions are reusable in other project
 */
```

```
#include <stdio.h>
#include <stdlib.h>
#include <assert.h>
```

```
#define BLOCK_SIZE 16
```

```
/*
*****
```

```
function name: gpu_matrix_mult
description: dot product of two matrix (not only square)
parameters:
    &a GPU device pointer to a m X n matrix (A)
    &b GPU device pointer to a n X k matrix (B)
    &c GPU device output purpose pointer to a m X k matrix (C)
    to store the result
```

```
Note:
```

```
grid and block should be configured as:
    dim3 dimGrid((k + BLOCK_SIZE - 1) / BLOCK_SIZE, (m + BLOCK_SIZE - 1) / BLOCK_SIZE);
    dim3 dimBlock(BLOCK_SIZE, BLOCK_SIZE);
```

```
further speedup can be obtained by using shared memory to decrease global memory access times
```

```
return: none
```

```
*****
```

```
*/
__global__ void gpu_matrix_mult(int *a,int *b, int *c, int m, int n, int k)
```

```
{
    int row = blockIdx.y * blockDim.y + threadIdx.y;
    int col = blockIdx.x * blockDim.x + threadIdx.x;
    int sum = 0;
    if( col < k && row < m)
    {
        for(int i = 0; i < n; i++)
        {
            sum += a[row * n + i] * b[i * k + col];
        }
        c[row * k + col] = sum;
    }
}
```

```
/*
*****
```

```
function name: gpu_square_matrix_mult
description: dot product of two matrix (not only square) in GPU
parameters:
```

```
    &a GPU device pointer to a n X n matrix (A)
    &b GPU device pointer to a n X n matrix (B)
```

```
&c GPU device output purpose pointer to a n X n matrix (C)
to store the result
```

Note:

grid and block should be configured as:

```
dim3 dim_grid((n - 1) / BLOCK_SIZE + 1, (n - 1) / BLOCK_SIZE + 1, 1);
dim3 dim_block(BLOCK_SIZE, BLOCK_SIZE, 1);
```

return: none

```
*****
```

```
*/
```

```
__global__ void gpu_square_matrix_mult(int *d_a, int *d_b, int *d_result, int n)
{
```

```
    __shared__ int tile_a[BLOCK_SIZE][BLOCK_SIZE];
    __shared__ int tile_b[BLOCK_SIZE][BLOCK_SIZE];
```

```
    int row = blockIdx.y * BLOCK_SIZE + threadIdx.y;
    int col = blockIdx.x * BLOCK_SIZE + threadIdx.x;
    int tmp = 0;
    int idx;
```

```
    for (int sub = 0; sub < gridDim.x; ++sub)
    {
```

```
        idx = row * n + sub * BLOCK_SIZE + threadIdx.x;
        if(idx >= n*n)
        {
            // n may not divisible by BLOCK_SIZE
            tile_a[threadIdx.y][threadIdx.x] = 0;
        }
        else
        {
            tile_a[threadIdx.y][threadIdx.x] = d_a[idx];
        }
    }
```

```
    idx = (sub * BLOCK_SIZE + threadIdx.y) * n + col;
    if(idx >= n*n)
    {
        tile_b[threadIdx.y][threadIdx.x] = 0;
    }
    else
    {
        tile_b[threadIdx.y][threadIdx.x] = d_b[idx];
    }
    __syncthreads();
```

```
    for (int k = 0; k < BLOCK_SIZE; ++k)
    {
        tmp += tile_a[threadIdx.y][k] * tile_b[k][threadIdx.x];
    }
    __syncthreads();
}
```

```
if(row < n && col < n)
{
    d_result[row * n + col] = tmp;
}
}
```

```
/*
```

```
*****
```

function name: gpu\_matrix\_transpose

description: matrix transpose

parameters:

```
&mat_in GPU device pointer to a rows X cols matrix
&mat_out GPU device output purpose pointer to a cols X rows matrix
to store the result
```

Note:

grid and block should be configured as:

```
dim3 dim_grid((n - 1) / BLOCK_SIZE + 1, (n - 1) / BLOCK_SIZE + 1, 1);
dim3 dim_block(BLOCK_SIZE, BLOCK_SIZE, 1);
```

return: none

```
*****
```

```
*/
```

```
__global__ void gpu_matrix_transpose(int* mat_in, int* mat_out, unsigned int rows, unsigned int cols)
{
```

```
    unsigned int idx = blockIdx.x * blockDim.x + threadIdx.x;
    unsigned int idy = blockIdx.y * blockDim.y + threadIdx.y;
```

```
    if (idx < cols && idy < rows)
    {
```

```
        unsigned int pos = idy * cols + idx;
```

```

        unsigned int pos = idx * cols + idy;
        unsigned int trans_pos = idx * rows + idy;
        mat_out[trans_pos] = mat_in[pos];
    }
}
/*
*****
function name: cpu_matrix_mult
description: dot product of two matrix (not only square) in CPU,
            for validating GPU results
parameters:
            &a CPU host pointer to a m X n matrix (A)
            &b CPU host pointer to a n X k matrix (B)
            &c CPU host output purpose pointer to a m X k matrix (C)
            to store the result
return: none
*****
*/
void cpu_matrix_mult(int *h_a, int *h_b, int *h_result, int m, int n, int k) {
    for (int i = 0; i < m; ++i)
    {
        for (int j = 0; j < k; ++j)
        {
            int tmp = 0.0;
            for (int h = 0; h < n; ++h)
            {
                tmp += h_a[i * n + h] * h_b[h * k + j];
            }
            h_result[i * k + j] = tmp;
        }
    }
}

/*
*****
function name: main
description: test and compare
parameters:
            none
return: none
*****
*/
int main(int argc, char const *argv[])
{
    int m, n, k;
    /* Fixed seed for illustration */
    srand(3333);
    printf("please type in m n and k\n");
    scanf("%d %d %d", &m, &n, &k);

    // allocate memory in host RAM, h_cc is used to store CPU result
    int *h_a, *h_b, *h_c, *h_cc;
    cudaMallocHost((void **) &h_a, sizeof(int)*m*n);
    cudaMallocHost((void **) &h_b, sizeof(int)*n*k);
    cudaMallocHost((void **) &h_c, sizeof(int)*m*k);
    cudaMallocHost((void **) &h_cc, sizeof(int)*m*k);

    // random initialize matrix A
    for (int i = 0; i < m; ++i) {
        for (int j = 0; j < n; ++j) {
            h_a[i * n + j] = rand() % 1024;
        }
    }

    // random initialize matrix B
    for (int i = 0; i < n; ++i) {
        for (int j = 0; j < k; ++j) {
            h_b[i * k + j] = rand() % 1024;
        }
    }

    float gpu_elapsed_time_ms, cpu_elapsed_time_ms;

    // some events to count the execution time
    cudaEvent_t start, stop;
    cudaEventCreate(&start);
    cudaEventCreate(&stop);

```

```

// start to count execution time of GPU version
cudaEventRecord(start, 0);
// Allocate memory space on the device
int *d_a, *d_b, *d_c;
cudaMalloc((void **) &d_a, sizeof(int)*m*n);
cudaMalloc((void **) &d_b, sizeof(int)*n*k);
cudaMalloc((void **) &d_c, sizeof(int)*m*k);

// copy matrix A and B from host to device memory
cudaMemcpy(d_a, h_a, sizeof(int)*m*n, cudaMemcpyHostToDevice);
cudaMemcpy(d_b, h_b, sizeof(int)*n*k, cudaMemcpyHostToDevice);

unsigned int grid_rows = (m + BLOCK_SIZE - 1) / BLOCK_SIZE;
unsigned int grid_cols = (k + BLOCK_SIZE - 1) / BLOCK_SIZE;
dim3 dimGrid(grid_cols, grid_rows);
dim3 dimBlock(BLOCK_SIZE, BLOCK_SIZE);

// Launch kernel
if(m == n && n == k)
{
    gpu_square_matrix_mult<<<dimGrid, dimBlock>>>(d_a, d_b, d_c, n);
}
else
{
    gpu_matrix_mult<<<dimGrid, dimBlock>>>(d_a, d_b, d_c, m, n, k);
}
// Transefr results from device to host
cudaMemcpy(h_c, d_c, sizeof(int)*m*k, cudaMemcpyDeviceToHost);
cudaThreadSynchronize();
// time counting terminate
cudaEventRecord(stop, 0);
cudaEventSynchronize(stop);

// compute time elapse on GPU computing
cudaEventElapsedTime(&gpu_elapsed_time_ms, start, stop);
printf("Time elapsed on matrix multiplication of %dx%d . %dx%d on GPU: %f ms.\n\n", m, n, n, k, gpu_elapsed_time_ms);

// start the CPU version
cudaEventRecord(start, 0);

cpu_matrix_mult(h_a, h_b, h_cc, m, n, k);

cudaEventRecord(stop, 0);
cudaEventSynchronize(stop);
cudaEventElapsedTime(&cpu_elapsed_time_ms, start, stop);
printf("Time elapsed on matrix multiplication of %dx%d . %dx%d on CPU: %f ms.\n\n", m, n, n, k, cpu_elapsed_time_ms);

// validate results computed by GPU
int all_ok = 1;
for (int i = 0; i < m; ++i)
{
    for (int j = 0; j < k; ++j)
    {
        //printf("[%d][%d]:%d == [%d][%d]:%d, ", i, j, h_cc[i*k + j], i, j, h_c[i*k + j]);
        if(h_cc[i*k + j] != h_c[i*k + j])
        {
            all_ok = 0;
        }
    }
    //printf("\n");
}

// roughly compute speedup
if(all_ok)
{
    printf("all results are correct!!!, speedup = %f\n", cpu_elapsed_time_ms / gpu_elapsed_time_ms);
}
else
{
    printf("incorrect results\n");
}

// free memory
cudaFree(d_a);
cudaFree(d_b);
cudaFree(d_c);
cudaFreeHost(h_a);

```

```
cudaFreeHost(h_b);  
cudaFreeHost(h_c);  
cudaFreeHost(h_cc);  
return 0;  
}
```

please type in m n and k

Time elapsed on matrix multiplication of -1638093008x21908 . 21908x-1666910242 on GPU: 1.264448 ms.

Time elapsed on matrix multiplication of -1638093008x21908 . 21908x-1666910242 on CPU: 0.002208 ms.

all results are correct!!!, speedup = 0.001746

[Colab paid products](#) - [Cancel contracts here](#)

✓ 1s completed at 5:47 PM

